



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

AY: 2026-27

Class:	TE-AI&DS	Semester:	V
Course Code:	2015111	Course Name:	Statistics for Machine Learning and Data Science Lab

Name of Student:	Bhavik Vijay Gharat
Roll No. :	69
Experiment No.:	5
Title of the Experiment:	Implement Resampling Techniques and Confidence Interval Estimation for Statistical Decision Making
Date of Performance:	
Date of Submission:	

Evaluation

Performance Indicator	Max. Marks	Marks Obtained
Performance	5	
Understanding	5	
Journal work and timely submission	10	
Total	20	

Performance Indicator	Exceed Expectations (EE)	Meet Expectations (ME)	Below Expectations (BE)
Performance	4-5	2-3	1
Understanding	4-5	2-3	1
Journal work and timely submission	8-10	5-8	1-4

Checked by

Name of Faculty :

Signature :

Date :

2015111: Statistics for Machine Learning and Data Science Lab



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Aim: To implement bootstrap and permutation resampling techniques on the Pima Indians Diabetes Dataset for estimating confidence intervals and analyzing statistical differences between groups.

Objectives

After completing this experiment, the student will be able to:

1. Apply bootstrap and permutation resampling techniques to estimate sampling distributions and confidence intervals.
2. Interpret resampling results to support statistical decision making.

Tools Required

- Python 3.x
- Google Colab / Jupyter Notebook
- Pandas
- NumPy
- SciPy
- Matplotlib
- Seaborn

Dataset Details

Dataset: Pima Indians Diabetes Dataset

The dataset contains diagnostic information for 768 female patients. The Outcome variable indicates whether the patient is diabetic.



Procedure

1. Load the Pima Indians Diabetes Dataset.
2. Select a numerical variable such as Glucose or BMI.
3. Draw repeated bootstrap samples from the original dataset with replacement.
4. Calculate the selected statistic, such as the mean, for every bootstrap sample.
5. Construct a bootstrap confidence interval using the obtained sampling distribution.
6. Divide the data into two groups based on Outcome.
7. Calculate the observed difference between the group means.
8. Randomly permute the group labels repeatedly.
9. Calculate the difference in group means for every permutation.
10. Compare the observed difference with the permutation distribution and interpret the result.

Description of the Experiment

Resampling methods repeatedly create new samples from the available data to study the behaviour of a statistic.

The experiment follows: Original Dataset, Repeated Resampling, Sampling Distribution, Confidence Interval / Statistical Comparison, and Interpretation

Two resampling techniques are implemented:

- Bootstrap: Used to estimate the variability and confidence interval of a statistic.



- Permutation Test: Used to test whether the observed difference between groups could have occurred by chance.

Detailed Description of Statistical Methods

Bootstrap Resampling

Bootstrap samples are created by repeatedly sampling observations with replacement from the original dataset.

For each bootstrap sample, a statistic such as the mean is calculated.

After many repetitions, the resulting values form an approximate sampling distribution.

A confidence interval can then be obtained from this distribution.

For example, a 95% percentile bootstrap confidence interval may be calculated using:

$$[Q_{.025}, Q_{.975}]$$

where $Q_{.025}$ and $Q_{.975}$ are the 2.5th and 97.5th percentiles of the bootstrap distribution.

Permutation Test

A permutation test evaluates whether the observed difference between two groups could have occurred by chance.

For example:

$$\text{Difference in Mean Glucose} = \bar{x}^{\text{diabetic}} - \bar{x}^{\text{non-diabetic}}$$

The group labels are randomly shuffled repeatedly while preserving the observations.

The resulting differences form the null distribution.



The p-value is estimated as the proportion of permuted differences that are at least as extreme as the observed difference.

Code:-

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("diabetes.csv")

print("Dataset Shape:", df.shape)
print("\nFirst 5 Records:")
print(df.head())

glucose = df["Glucose"].dropna()

print("\nOriginal Mean Glucose:", glucose.mean())

np.random.seed(42)

n_bootstrap = 10000
bootstrap_means = []

for i in range(n_bootstrap):
    sample = np.random.choice(
        glucose,
        size=len(glucose),
        replace=True
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

)

```
bootstrap_means.append(np.mean(sample))
```

```
bootstrap_means = np.array(bootstrap_means)
```

```
lower = np.percentile(bootstrap_means, 2.5)
```

```
upper = np.percentile(bootstrap_means, 97.5)
```

```
print("\nBOOTSTRAP RESULTS")
```

```
print("Original Mean Glucose:", glucose.mean())
```

```
print("95% Bootstrap Confidence Interval:")
```

```
print("Lower Limit:", lower)
```

```
print("Upper Limit:", upper)
```

```
plt.figure(figsize=(8, 5))
```

```
sns.histplot(bootstrap_means, bins=40, kde=True)
```

```
plt.axvline(lower, linestyle="--", label="2.5th Percentile")
```

```
plt.axvline(upper, linestyle="--", label="97.5th Percentile")
```

```
plt.axvline(glucose.mean(), linestyle="-", label="Original Mean")
```

```
plt.title("Bootstrap Distribution of Mean Glucose")
```

```
plt.xlabel("Mean Glucose")
```

```
plt.ylabel("Frequency")
```

```
plt.legend()
```

```
plt.show()
```

```
diabetic = df[df["Outcome"] == 1]["Glucose"].dropna()
```

```
non_diabetic = df[df["Outcome"] == 0]["Glucose"].dropna()
```



```
observed_difference = diabetic.mean() - non_diabetic.mean()

print("\nPERMUTATION TEST RESULTS")
print("Mean Glucose - Diabetic:", diabetic.mean())
print("Mean Glucose - Non-Diabetic:", non_diabetic.mean())
print("Observed Difference:", observed_difference)

combined = np.concatenate([diabetic.values, non_diabetic.values])

n_diabetic = len(diabetic)
n_permutations = 10000

permutation_differences = []

for i in range(n_permutations):
    shuffled = np.random.permutation(combined)

    perm_diabetic = shuffled[:n_diabetic]
    perm_non_diabetic = shuffled[n_diabetic:]

    difference = np.mean(perm_diabetic) - np.mean(perm_non_diabetic)
    permutation_differences.append(difference)

permutation_differences = np.array(permutation_differences)

p_value = np.mean(
    np.abs(permutation_differences) >= abs(observed_difference)
)
```



```
print("Permutation p-value:", p_value)

plt.figure(figsize=(8, 5))

sns.histplot(permutation_differences, bins=40, kde=True)

plt.axvline(
    observed_difference,
    linestyle="--",
    label="Observed Difference"
)

plt.axvline(
    -observed_difference,
    linestyle="--"
)

plt.title("Permutation Distribution of Difference in Mean Glucose")
plt.xlabel("Difference in Mean Glucose")
plt.ylabel("Frequency")
plt.legend()
plt.show()

alpha = 0.05


print("\nSTATISTICAL DECISION")


if p_value < alpha:
    print("Reject the null hypothesis.")
```



```
print("There is a statistically significant difference")
print("in mean glucose levels between the two groups.")
else:
    print("Fail to reject the null hypothesis.")
    print("There is no statistically significant difference")
    print("in mean glucose levels between the two groups.")
```

Output:-





```
Dataset Shape: (768, 9)

First 5 Records:
  Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunction  Age  Outcome
0           6     148             72           35         0   33.6              0.627   50         1
1           1      85             66           29         0   26.6              0.351   31         0
2           8     183             64           0          0   23.3              0.672   32         1
3           1      89             66           23        94   28.1              0.167   21         0
4           0     137             40           35        168  43.1              2.288   33         1

Original Mean Glucose: 120.89453125

BOOTSTRAP RESULTS
Original Mean Glucose: 120.89453125
95% Bootstrap Confidence Interval:
Lower Limit: 118.69401041666667
Upper Limit: 123.11588541666667

PERMUTATION TEST RESULTS
Mean Glucose - Diabetic: 141.25746268656715
Mean Glucose - Non-Diabetic: 109.98
Observed Difference: 31.277462686567148
Permutation p-value: 0.0

STATISTICAL DECISION
Reject the null hypothesis.
There is a statistically significant difference
in mean glucose levels between the two groups.
```



Conclusion

Bootstrap and permutation resampling techniques were implemented on the Pima Indians Diabetes Dataset. Bootstrap resampling was used to estimate the sampling distribution and confidence interval of a selected statistic, while permutation testing was used to assess the significance of differences between groups.

Questions to be Answered by Students

1.What statistic was selected for bootstrap resampling, and what confidence interval was obtained?

The mean Glucose level was selected for bootstrap resampling. A 95% bootstrap confidence interval was obtained using the 2.5th and 97.5th percentiles of the bootstrap distribution.

2.How does the bootstrap distribution compare with the original sample statistic?

The bootstrap distribution is centered close to the original sample mean. This indicates that repeated resampling gives similar estimates of the mean and shows that the bootstrap estimate is stable.

3.What was the observed difference between the selected groups?

The observed difference was the difference between the mean Glucose levels of diabetic and non-diabetic patients. It was calculated as:

Mean Glucose (Diabetic) – Mean Glucose (Non-diabetic).

4.Based on the permutation test, is there sufficient statistical evidence to conclude that the observed group difference is significant?

Yes. Since the permutation p-value is expected to be less than 0.05, we reject the null hypothesis. Therefore, there is sufficient statistical evidence that the mean Glucose levels differ significantly between diabetic and non-diabetic groups.

5.What effect would increasing the number of bootstrap or permutation repetitions have on the results?

Increasing the number of repetitions makes the estimated confidence interval and p-value more stable and accurate. It reduces random variation caused by resampling, although it also increases computation time.